

# Consistency and Replication

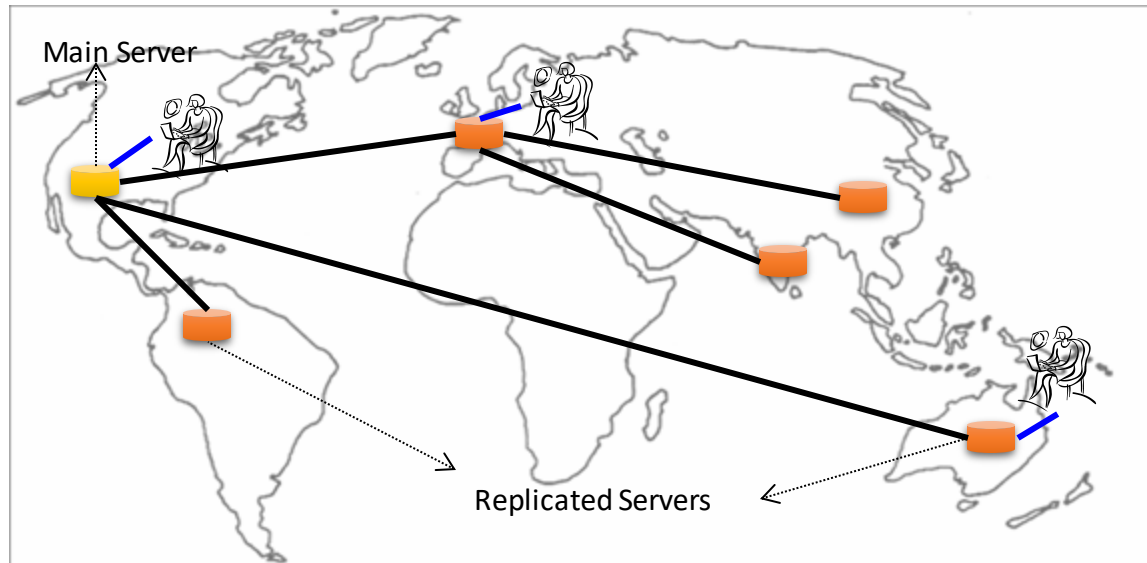
# Why Replication?

- Replication is the process of maintaining the data at multiple computers
- Replication is necessary for:
  1. Improving performance
    - A client can access the replicated copy of the data that is near to its location
  2. Increasing the availability of services
    - Replication can mask failures such as server crashes and network disconnection
  3. Enhancing the scalability of the system
    - Requests to the data can be distributed to many servers which contain replicated copies of the data
  4. Securing against malicious attacks
    - Even if some replicas are malicious, secure data can be guaranteed to the client by relying on the replicated copies at the non-compromised servers

# 1. Replication for Improving Performance

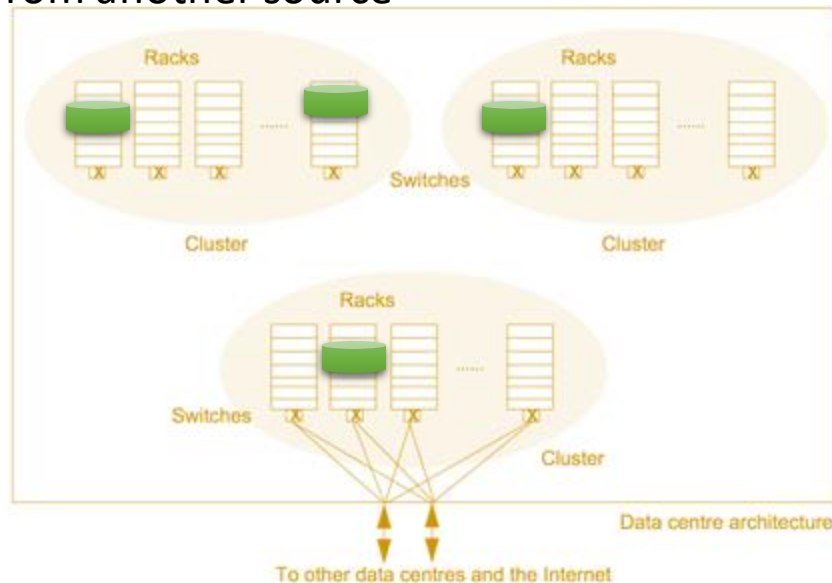
- Example Applications

- Caching webpages at the client browser
- Caching IP addresses at clients and DNS Name Servers
- Caching in Content Delivery Network (CDNs)
  - Commonly accessed contents, such as software and streaming media, are cached at various network locations



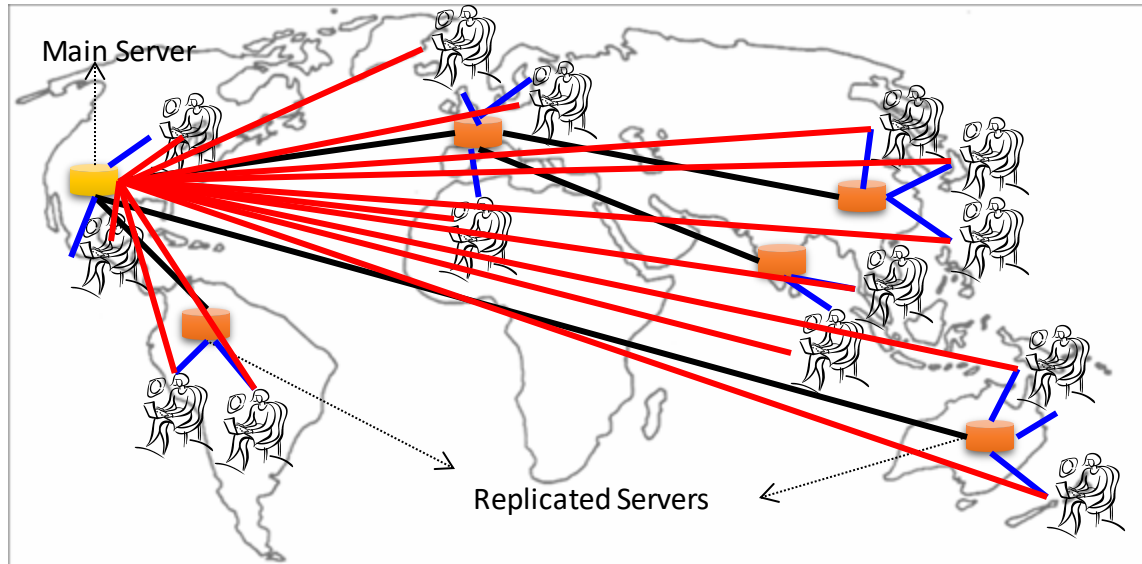
## 2. Replication for High-Availability

- Availability can be increased by storing the data at replicated locations (instead of storing one copy of the data at a server)
- Example: Google File-System and Chubby replicate the data at computers across different racks, clusters and data-centers
  - If one computer or a rack or a cluster crashes, then the data can still be accessed from another source




### 3. Replication for Enhancing Scalability

- Distributing the data across replicated servers helps in avoiding bottle-necks at the main server
  - It balances the load between the main and the replicated servers
- Example: Content Delivery Networks decrease the load on main servers of the website



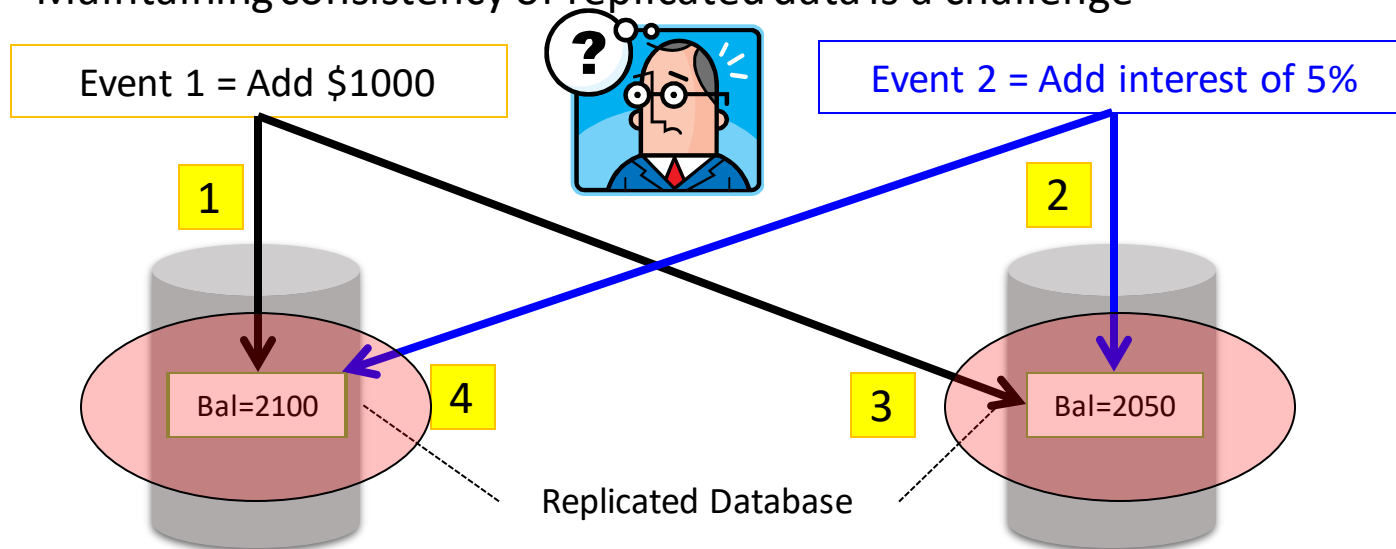
## 4. Replication for Securing Against Malicious Attacks

- If a minority of the servers that hold the data are malicious, the non-malicious servers can outvote the malicious servers, thus providing security.
- The technique can also be used to provide fault-tolerance against non-malicious but faulty servers
- Example: In a peer-to-peer system, peers can coordinate to prevent delivering faulty data to the requester



# Why Consistency?

- In a DS with replicated data, one of the main problems is keeping the data consistent
- An example:
  - In an e-commerce application, the bank database has been replicated across two servers
  - Maintaining consistency of replicated data is a challenge



# Overview of Consistency and Replication

## TODAY

- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
  - When, where and by whom replicas should be placed?
  - Which consistency model to use for keeping replicas consistent?
- Consistency Protocols
  - We study various implementations of consistency models

## NEXT




# Overview

- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

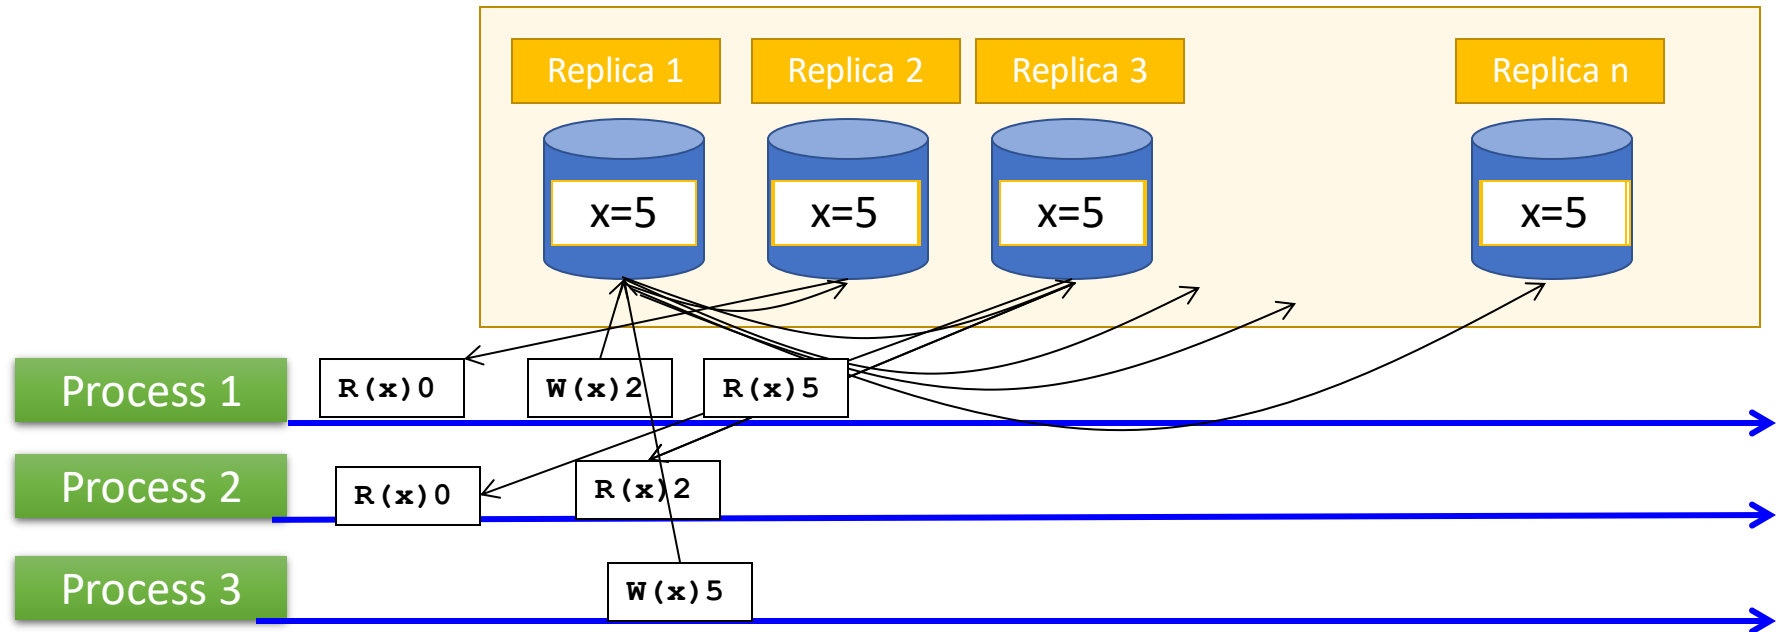
# Introduction to Consistency and Replication

- In a distributed system, shared data is typically stored in distributed shared memory, distributed databases or distributed file systems.
  - The storage can be distributed across multiple computers
  - Simply, we refer to a series of such data storage units as *data-stores*
- Multiple processes can access shared data by accessing any replica on the data-store
  - Processes generally perform read and write operations on the replicas



# Maintaining Consistency of Replicated Data

## DATA-STORE



## Strict Consistency

- Data is always fresh
  - After a write operation, the update is propagated to all the replicas
  - A read operation will result in reading the most recent write
- If there are occasional writes and reads, this leads to large overheads

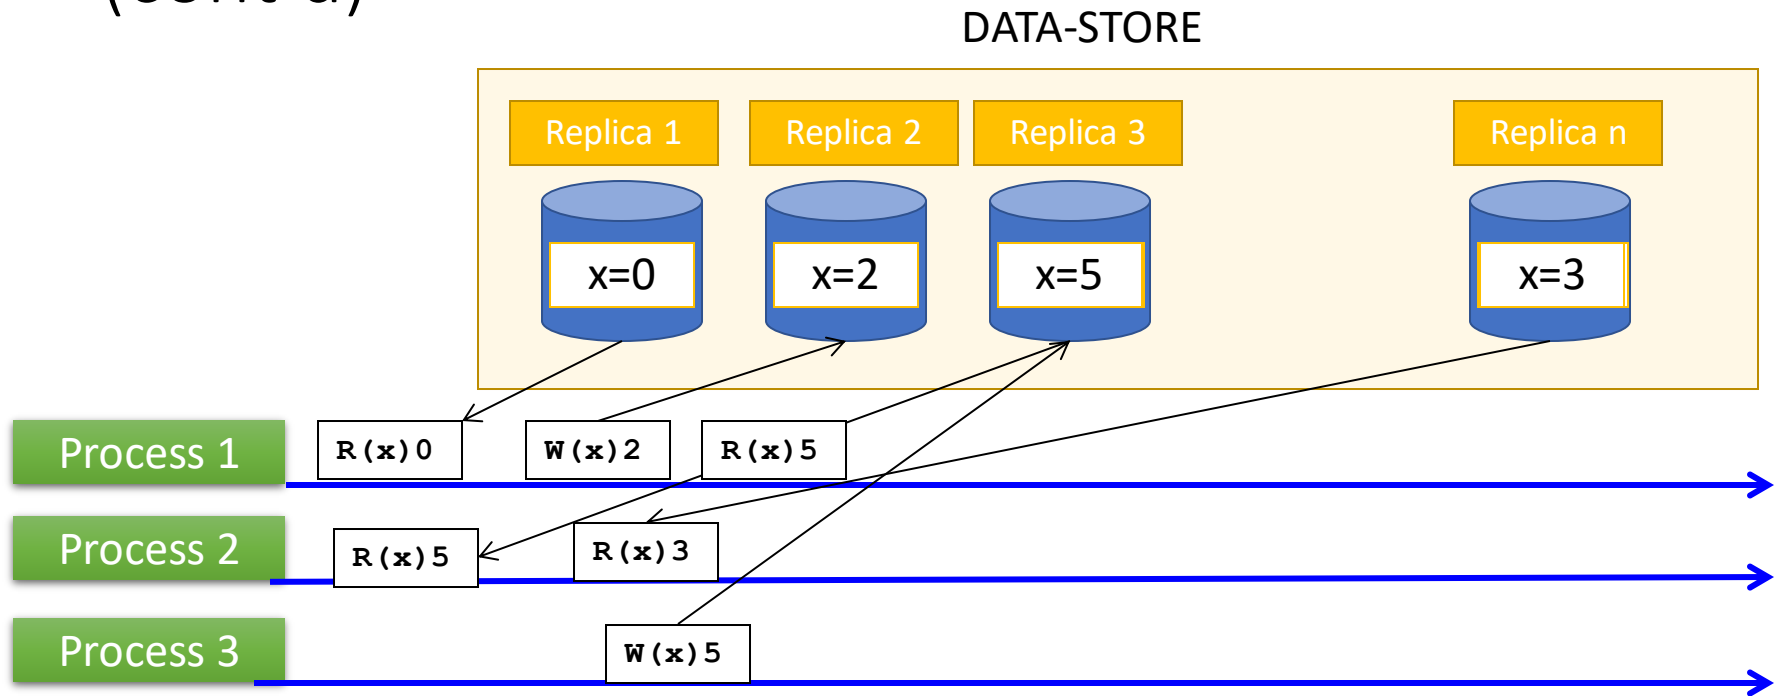
**P1** = Process P1

$\longrightarrow$  = Timeline at P1

**$R(x)b$**  = Read variable  $x$ ;  
Result is  $b$

**$W(x)b$**  = Write variable  $x$ ;  
Result is  $b$

# Maintaining Consistency of Replicated Data (cont'd)



## Loose Consistency

- Data might be stale
  - A read operation may result in reading a value that was written long back
  - Replicas are generally out-of-sync
- The replicas may sync at coarse grained time, thus reducing the overhead

**P1** = Process P1


$\longrightarrow$  = Timeline at P1

$R(x)b$  = Read variable  $x$ ;  
Result is  $b$

$W(x)b$  = Write variable  $x$ ;  
Result is  $b$

# Trade-offs in Maintaining Consistency

- Maintaining consistency should balance between the strictness of consistency versus efficiency
  - Good-enough consistency depends on your application



# Consistency Model

- A consistency model is a contract between
  - the process that wants to use the data, and
  - the replicated data repository (or data-store)
- A consistency model states the level of consistency provided by the *data-store* to the processes while reading and writing the data

# Types of Consistency Models

- Consistency models can be divided into two types:
  - Data-Centric Consistency Models
    - These models define how the data updates are propagated across the replicas to keep them consistent
  - Client-Centric Consistency Models
    - These models assume that clients connect to different replicas at each time
    - The models ensure that whenever a client connects to a replica, the replica is brought up to date with the replica that the client accessed previously

# Overview

- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols



# Data-centric Consistency Models

- Data-centric Consistency Models describe how the replicated data is kept consistent, and what the process can expect
- Under Data-centric Consistency Models, we study two types of models:
  - Consistency Specification Models:
    - These models enable specifying the consistency levels that are tolerable to the application
  - Models for Consistent Ordering of Operations:
    - These models specify the order in which the data updates are propagated to different replicas

# Overview

- Consistency Models
  - Data-Centric Consistency Models
    - Consistency Specification Models
    - Models for Consistent Ordering of Operations
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

# Consistency Specification Models

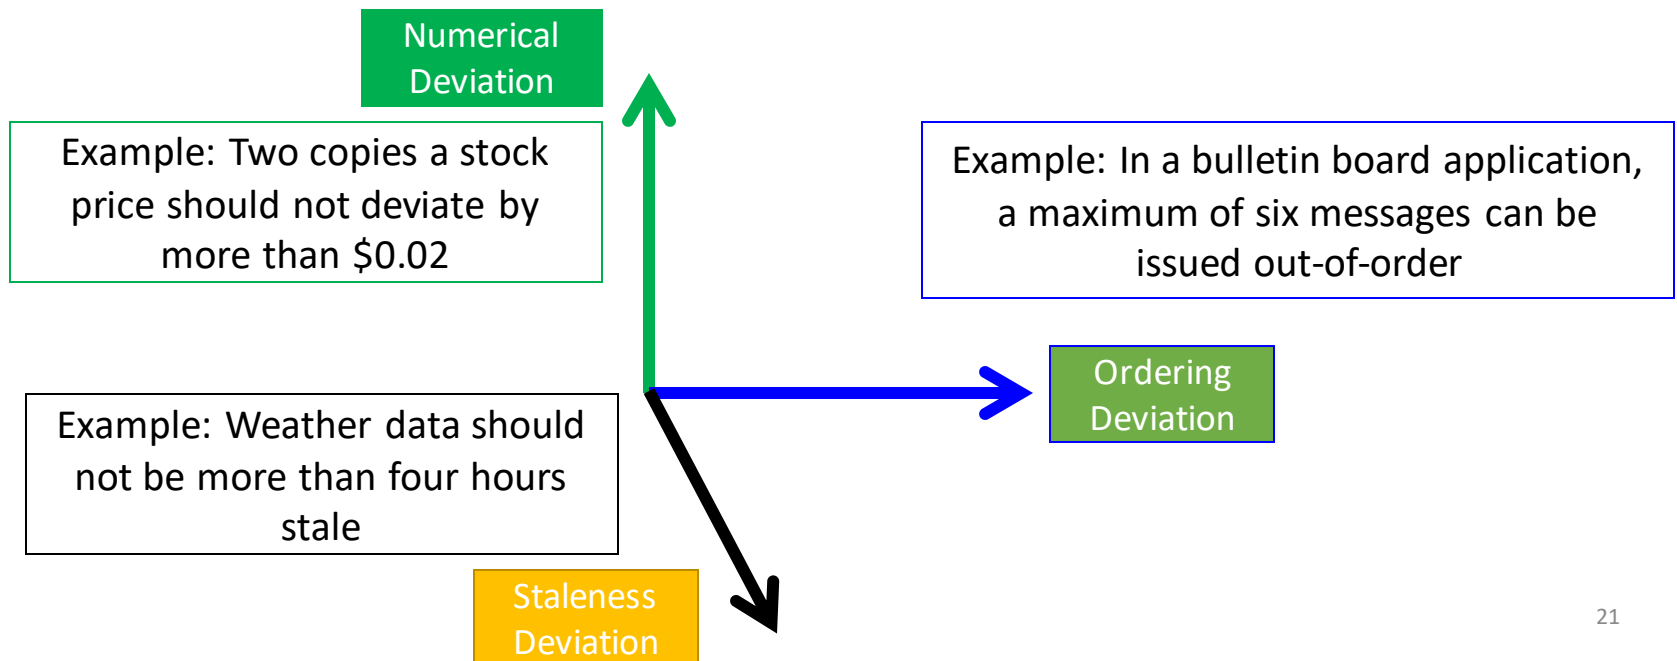
- In replicated data-stores, there should be a mechanism to:
  - Measure how inconsistent the data might be on different replicas
  - How replicas and applications can specify the tolerable inconsistency levels
- Consistency Specification Models enable measuring and specifying the level of inconsistency in a replicated data-store
- We study a Consistency Specification Model called *Continuous Consistency Model*

# Continuous Consistency Model

- Continuous Consistency Model is used to measure inconsistencies and express what inconsistencies can be expected in the system
- Yu and Vahdat [1] provided a framework for measuring and expressing consistency in replicated data-stores

# Continuous Consistency Ranges

- Level of consistency is defined over three independent axes:
  - **Numerical Deviation**: Deviation in the numerical values between replicas
  - **Order Deviation**: Deviation with respect to the ordering of update operations
  - **Staleness Deviation**: Deviation in the staleness between replicas



# Consistency Unit (Conit)

- Consistency unit (Conit) specifies the data unit over which consistency is measured
  - For example, conit can be defined as a record representing a single stock
- Level of consistency is measured by each replica along the three dimensions
  - Numerical Deviation
    - For a given replica R, how many updates at other replicas are not yet seen at R?  
What is the effect of the non-propagated updates on local Conit values?
  - Order Deviation
    - For a given replica R, how many local updates are not propagated to other replicas?
  - Staleness Deviation
    - For a given replica R, how long has it been since updates were propagated?

# Example of Conit and Consistency Measures

Order Deviation at a replica R is the number of operations in R that are not present at the other replicas

Numerical Deviation at replica R is defined as  $n(w)$ , where  
 $n$  = # of operations at other replicas that are not yet seen by R,  
 $w$  = weight of the deviation  
 = max(update amount of all variables in a Conit)

Replica A					Replica B				
x	y	VC	Ord	Num	x	y	VC	Ord	Num
0	0	(0,0)	0	0(0)	0	0	(0,0)	0	0(0)
0	0	(0,0)	0	1(2)	2	0	(0,5)	1	0(0)
2	0	(1,5)	0	0(0)	2	0	(0,5)	0	0(0)
2	1	(10,5)	1	0(0)	2	0	(0,5)	0	1(1)
2	1	(10,5)	1	1(1)	2	1	(0,16)	1	1(1)
3	1	(14,5)	2	1(1)	2	1	(0,16)	1	2(2)
3	4	(23,5)	3	1(1)	2	1	(0,16)	1	3(4)

## Replica A

x; y

Operation		Result
<5,B>	x+=2	x=2
<10,A>	y+=1	y=1
<14,A>	x+=1	x=3
<23,A>	y+=3	y=4

## Replica B

x; y

Operation		Result
<5,B>	x+=2	x=2
<16,B>	y+=1	y=1

<5,B> =

Operation performed at B  
when the vector clock was 5

<m,n>

= Uncommitted  
operation

<m,n>

= Committed  
operation

x;y

= A Conit

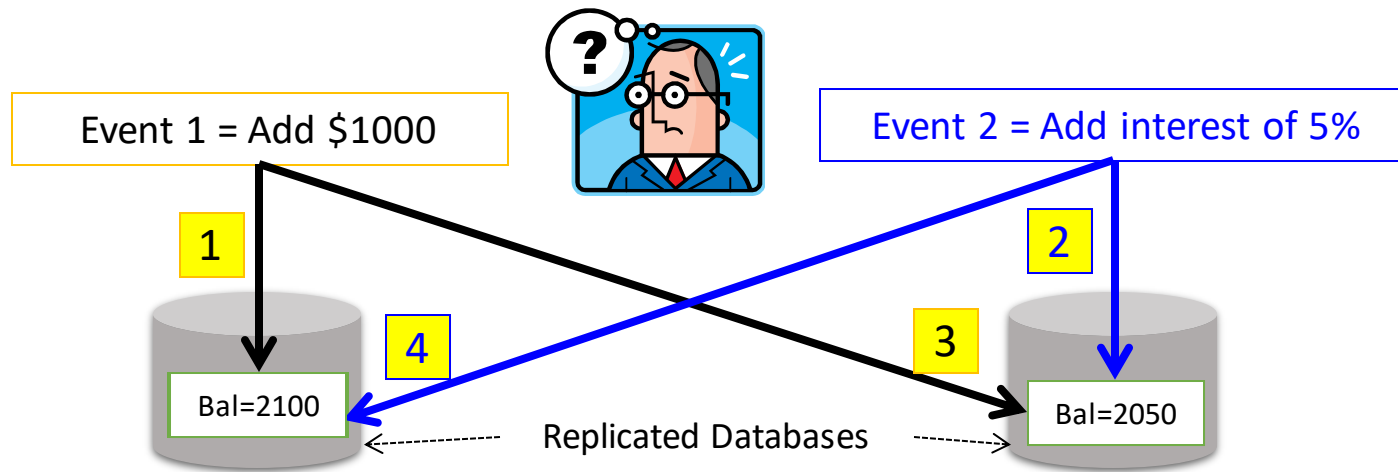
# Overview

- Consistency Models
  - Data-Centric Consistency Models
    - Continuous Specification Models
    - **Models for Consistent Ordering of Operations**
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols



# Why is Consistent Ordering Required in Replication?

- In several applications, the order or the sequence in which the replicas commit to the data store is critical
- Example:



- Continuous Specification Models defined how inconsistency is measured
  - However, the models did not enforce any order in which the data is committed

# Consistent Ordering of Operations (cont'd)

- Whenever a replica is updated, it propagates the updates to other replicas at some point in time
- Updating different replicas is carried out by passing messages between the replica data-stores
- We will study different types of ordering and consistency models arising from these orderings

# Types of Ordering

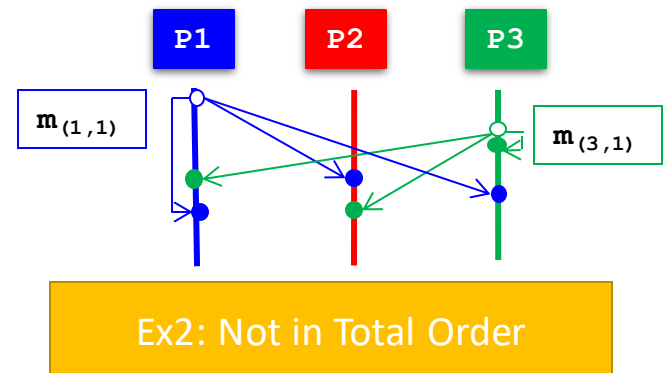

- We will study three types of ordering of messages that meet the needs of different applications:
  1. Total Ordering
  2. Sequential Ordering
    - i. Sequential Consistency Model
  3. Causal Ordering
    - i. Causal Consistency Model

# Types of Ordering

1. Total Ordering
2. Sequential Ordering
3. Causal Ordering

# Total Ordering

- Total Order
  - If process  $P_i$  sends a message  $m_i$  and  $P_j$  sends  $m_j$ , and if one correct process delivers  $m_i$  before  $m_j$  then every correct process delivers  $m_i$  before  $m_j$
- Messages can contain replica updates, such as passing the read or write operation that needs to be performed at each replica
  - In the example Ex1, if  $P_1$  issues the operation  $m_{(1,1)} : x=x+1;$  and
  - If  $P_3$  issues  $m_{(3,1)} : \text{print}(x);$
  - Then, at all replicas  $P_1, P_2, P_3$  the following order of operations are executed
    - $\text{print}(x);$
    - $x=x+1;$

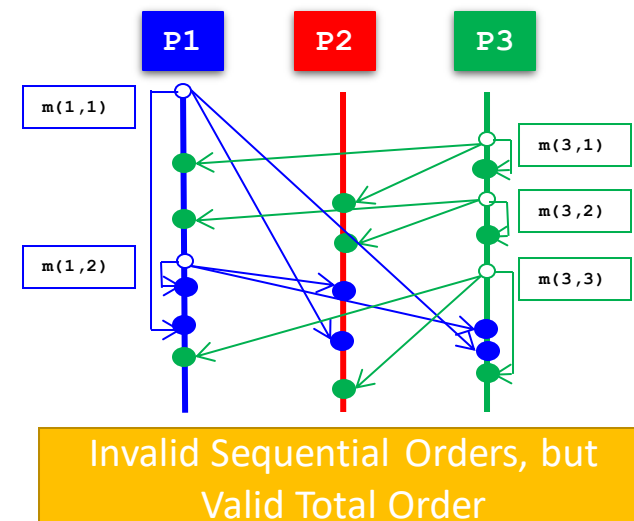
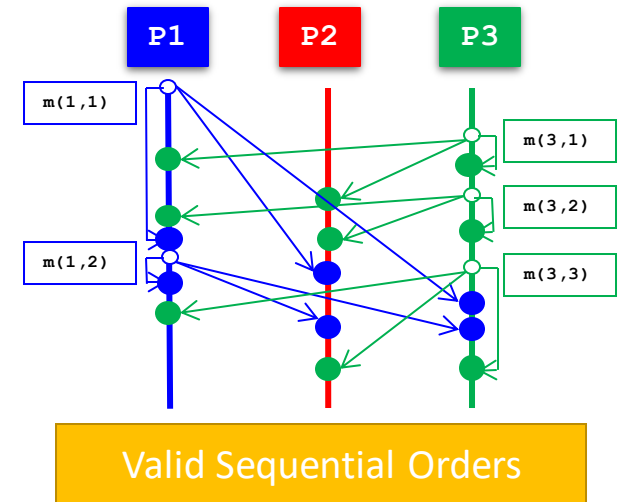


# Types of Ordering

1. Total Ordering
2. Sequential Ordering
3. Causal Ordering


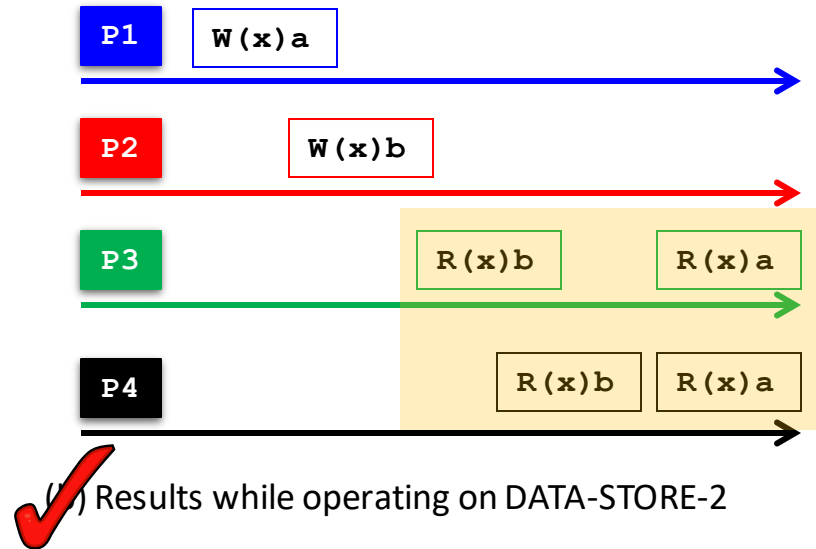
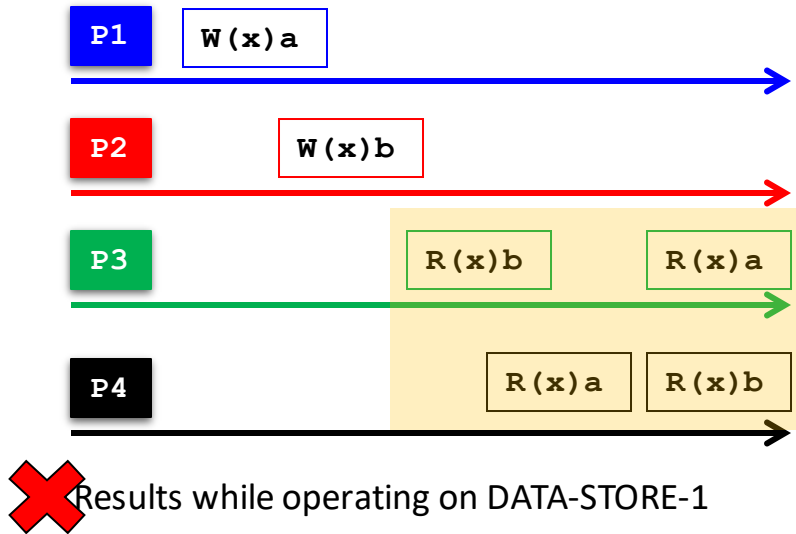
# Sequential Ordering

- + If a process  $P_i$  sends a sequence of messages  $m_{(i,1)}, \dots, m_{(i,n_i)}$ , and
- + Process  $P_j$  sends a sequence of messages  $m_{(j,1)}, \dots, m_{(j,n_j)}$ ,
- + Then, :
  - + At any process, the set of messages received are in some sequential order
  - + Messages from each individual process appear in this sequence in the order sent by the sender
    - + At every process,  $m_{i,1}$  should be delivered before  $m_{i,2}$ , which is delivered before  $m_{i,3}$  and so on...
    - + At every process,  $m_{j,1}$  should be delivered before  $m_{j,2}$ , which is delivered before  $m_{j,3}$  and so on...



# Sequential Consistency Model


- Sequential Consistency Model enforces that all the update operations are executed at the replicas in a sequential order
- Consider a data-store with variable **x** (Initialized to **NULL**)
  - In the two data-stores below, identify the sequentially consistent data-store





# Sequential Consistency (cont'd)

- Consider three processes  $P_1$ ,  $P_2$  and  $P_3$  executing multiple instructions on three shared variables  $x$ ,  $y$  and  $z$ 
  - Assume that  $x$ ,  $y$  and  $z$  are set to zero at start



- There are many valid sequences in which operations can be executed at the replica respecting sequential consistency
  - Identify the output

```

x = 1
print (y,z)
y = 1
print (x,z)
z = 1
print (x,y)
    
```

```

x = 1
y = 1
print (x,z)
print (y,z)
z = 1
print (x,y)
    
```

```

z = 1
print (x,y)
print (x,z)
y = 1
x = 1
print (y,z)
    
```

```

y = 1
z = 1
print (x,y)
print (x,z)
x = 1
print (y,z)
    
```

Output

001011

101011

000111

010111



# Implications of Adopting Sequential Consistency Model for Applications

- There might be several different sequentially consistent combinations of ordering
  - Number of combinations for a total of  $n$  instructions =  $O(n!)$
- The contract between the process and the distributed data-store is that the process must accept all of the sequential orderings as valid results
  - A process that works for some of the sequential orderings and does not work correctly for others is INCORRECT

# Summary

- Replication is necessary for improving performance, scalability and availability, and for providing fault-tolerance
- Replicated data-stores should be designed after carefully evaluating the trade-off between tolerable data inconsistency and efficiency
- Consistency Models describe the contract between the data-store and process about what form of consistency to expect from the system
- Data-centric consistency models:
  - Continuous Consistency Models provide mechanisms to measure and specify inconsistencies
  - Consistency Models can be defined based on the type of ordering of operations that the replica guarantees the applications
    - We studied Sequential Consistency Model

# Next ...

- Consistency Models
  - Data-Centric Consistency Model: Causal Consistency Model
  - Client-Centric Consistency Models
- Replica Management
  - Replica management studies:
    - when, where and by whom replicas should be placed
    - which consistency model to use for keeping replicas consistent
- Consistency Protocols
  - We study various implementations of consistency models

# References

- [\[1\] Haifeng Yu and Amin Vahdat, “Design and evaluation of a conit-based continuous consistency model for replicated services”](#)
- [\[2\] <http://tech.amikelive.com/node-285/using-content-delivery-networks-cdn-to-speed-up-content-load-on-the-web/>](#)
- [\[3\] \[http://en.wikipedia.org/wiki/Replication\\\_\\(computer\\\_science\\)\]\(http://en.wikipedia.org/wiki/Replication\_\(computer\_science\)\)](#)
- [\[4\] \[http://en.wikipedia.org/wiki/Content\\\_delivery\\\_network\]\(http://en.wikipedia.org/wiki/Content\_delivery\_network\)](#)
- [\[5\] <http://www.cdk5.net>](#)
- [\[6\] <http://www.dis.uniroma1.it/~baldoni/ordered%2520communication%25202008.ppt>](#)
- [\[7\] <http://www.cs.uiuc.edu/class/fa09/cs425/L5tmp.ppt>](#)